

Software Transactional Memory (5)

Mohamed Mohamedin

Chapter 4 of TM Book

Lock-based STMs with Global Metadata

- RingSTM
- Another STM with Global Metadata, BUT
 - Doesn't use value-based validation
 - Instead, it uses signature-based validation

RingSTM

- Bloom Filters
 - It is an array of bits
 - Represents a Set
 - It is a probabilistic data structure
 - It can tell if an element is possibly in the Set
 - Has False Positives
 - BUT, it can 100% tell that an element is NOT in the Set
 - Contains return
 - Possibly in the Set
 - Definitely NOT in the Set
 - Has only Add & Contains. No Remove

RingSTM

- Array of bits of a given size m
 - Initially all bits are zeros
 - Each element is hashed using k hash functions
 - Each one map an element to a bit
 - Set those bits to ones

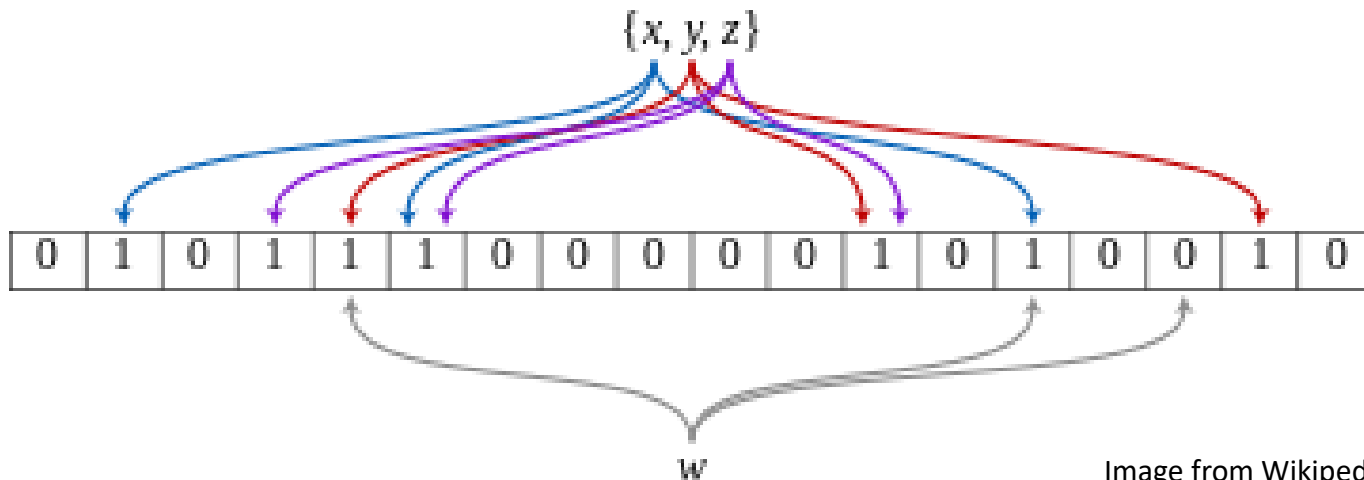


Image from Wikipedia

RingSTM

- Add(element)
 - Hash the element with each hash function
 - Set the corresponding bit to one
 - E.g., elements x , y , and z

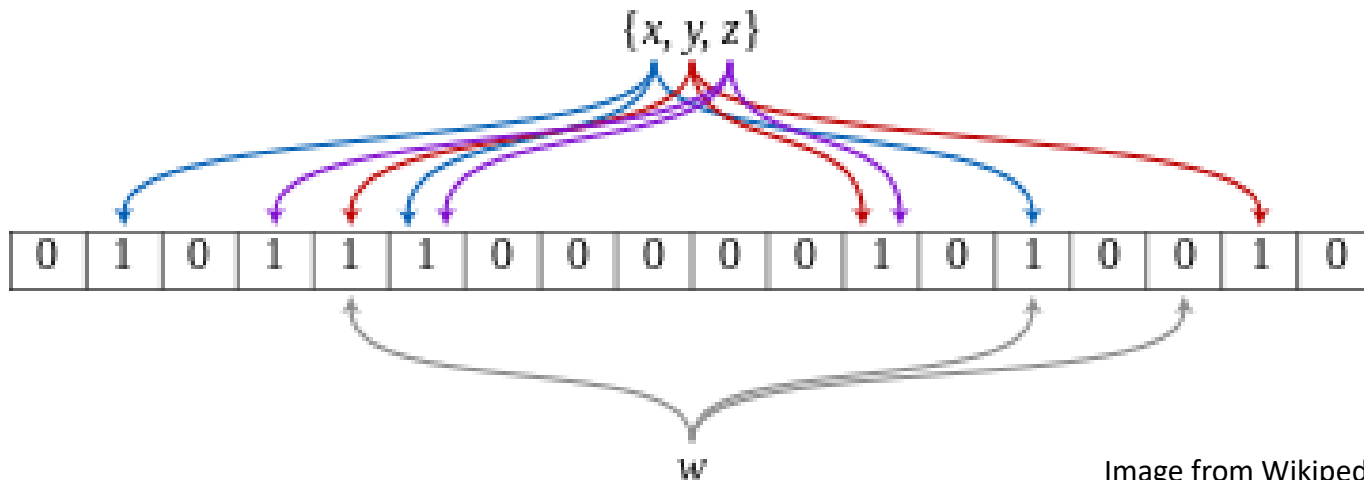


Image from Wikipedia

RingSTM

- Contains(element)
 - Hash the element with each hash function
 - If ALL corresponding bit are ones
 - Return true (possibly in the set)
 - Else: It is NOT in the set (definitely)
 - E.g., w is not in the Set




Image from Wikipedia

RingSTM

- Probability of false positives

- Adding too many elements to a small Bloom filter increased false-positives significantly

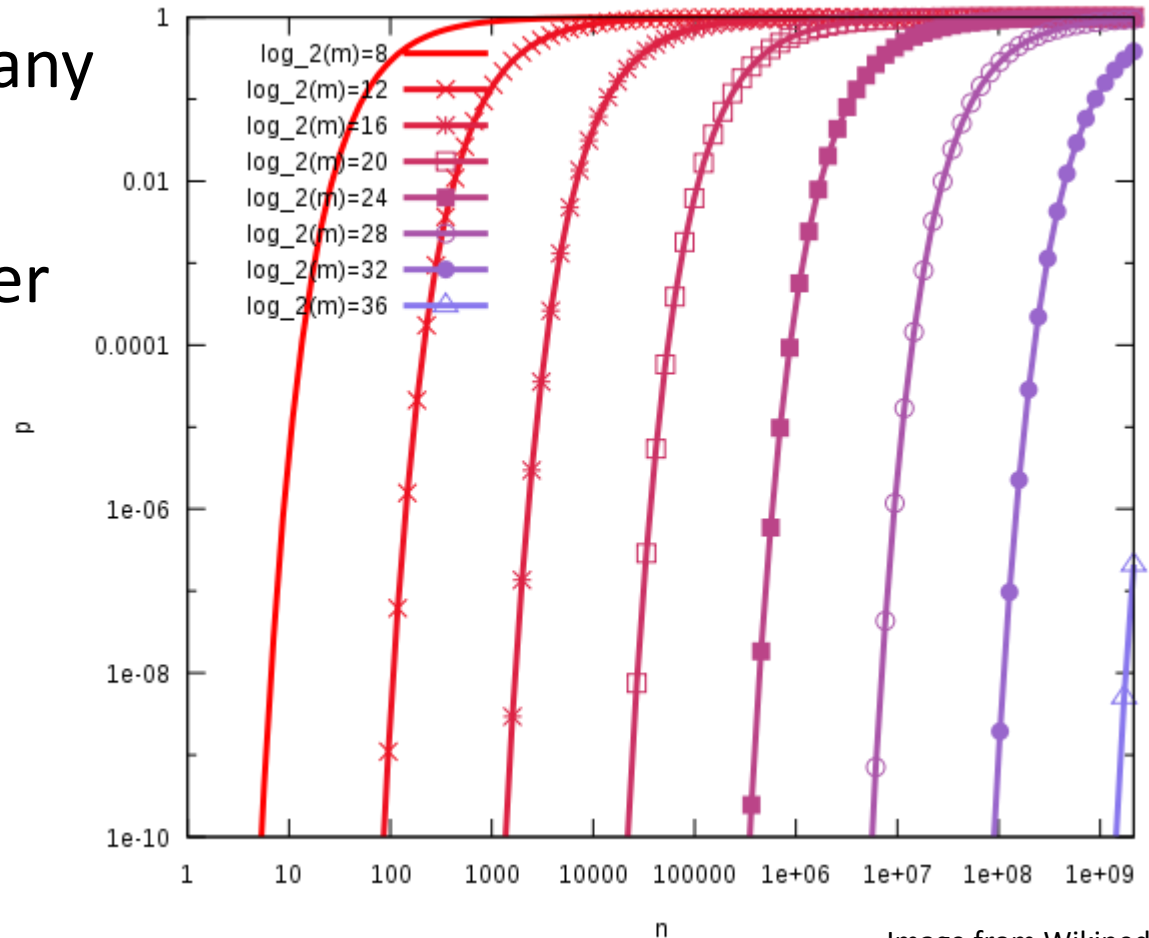


Image from Wikipedia

RingSTM

- Other methods can be defined
 - Intersect:
 - Check if two Bloom filters has common elements
 - Union:
 - Merge two sets (Bloom filters)

RingSTM

- A simplified Bloom filter implementation
 - From RSTM (<http://cs.rochester.edu/research/synchronization/rstm/>)
 - One hash function is used

RingSTM

```
template <uint32_t BITS>
class BitFilter
{
    static const uint32_t WORD_SIZE    = 8 * sizeof(uintptr_t);
    static const uint32_t WORD_BLOCKS = BITS / WORD_SIZE;

    uintptr_t word_filter[WORD_BLOCKS];

    static uint32_t hash(const void* const key)
    {
        return (((uintptr_t)key) >> 3) % BITS;
    }

public:
    BitFilter() { clear(); }

    void add(const void* const val) volatile
    {
        const uint32_t index = hash(val);
        const uint32_t block = index / WORD_SIZE;
        const uint32_t offset = index % WORD_SIZE;
        word_filter[block] |= (1u << offset);
    }
}
```

RingSTM

```
bool lookup(const void* const val) const volatile
{
    const uint32_t index = hash(val);
    const uint32_t block = index / WORD_SIZE;
    const uint32_t offset = index % WORD_SIZE;

    return word_filter[block] & (1u << offset);
}

void unionwith(const BitFilter<BITS>& rhs)
{
    for (uint32_t i = 0; i < WORD_BLOCKS; ++i)
        word_filter[i] |= rhs.word_filter[i];
}

void clear() volatile
{
    for (uint32_t i = 0; i < WORD_BLOCKS; ++i)
        word_filter[i] = 0;
}
```

RingSTM

```
void fastcopy(const volatile BitFilter<BITS>* rhs) volatile
{
    for (uint32_t i = 0; i < WORD_BLOCKS; ++i)
        word_filter[i] = rhs->word_filter[i];
}

bool intersect(const BitFilter<BITS>* rhs) const volatile
{
    for (uint32_t i = 0; i < WORD_BLOCKS; ++i)
        if (word_filter[i] & rhs->word_filter[i])
            return true;
    return false;
}
```

RingSTM

- How it works?
 - Metadata
 - Global:
 - Global-Clock (ring-index)
 - The Ring: “An ordered, fixed size ring data structure”
 - Thread-local:
 - Read-set signature
 - Write-set signature
 - Write-set
 - RV

RingSTM

- How it works?

- Metadata

- Global:

- Global-Clock (ring-in)
 - The Ring: “An ordered sequence of entries”

- Thread-local:

- Read-set signature
 - Write-set signature
 - Write-set
 - RV

Listing 1 RingSTM metadata


```
struct RingEntry
    int      ts           // commit timestamp
    filter   wf           // write filter
    int      prio        // priority
    int      st           // writing or complete

struct Transaction
    set      wset         // speculative writes
    filter   wf           // addresses to write
    filter   rf           // addresses read
    int      start        // logical start time

RingEntry ring[]         // the global ring
int        ring_index    // newest ring entry
int        prefix_index  // RingR prefix field
```

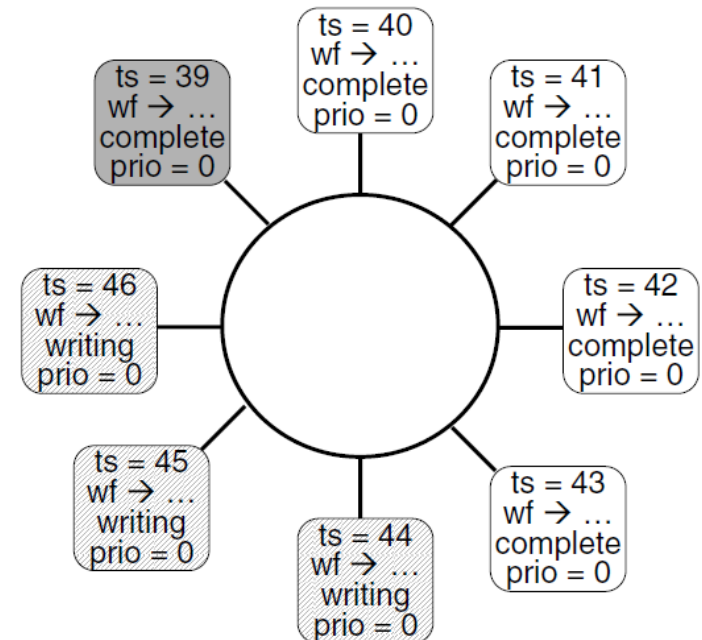
The Ring

- Circular data structure
- Hold committed transaction info
 - Commit Timestamp (ts)
 - Write-signature (wf)
 - Status
 - Priority
 - Initially:
 - All have ts = 0
 - Status = complete



The Ring

- Only write transactions modifies the Ring
 - One CAS operation to add an entry
 - A transaction is committed (logically) once its entry is added to the Ring (status: writing)
 - After writing back is finished
 - Status: complete
 - Physically committed

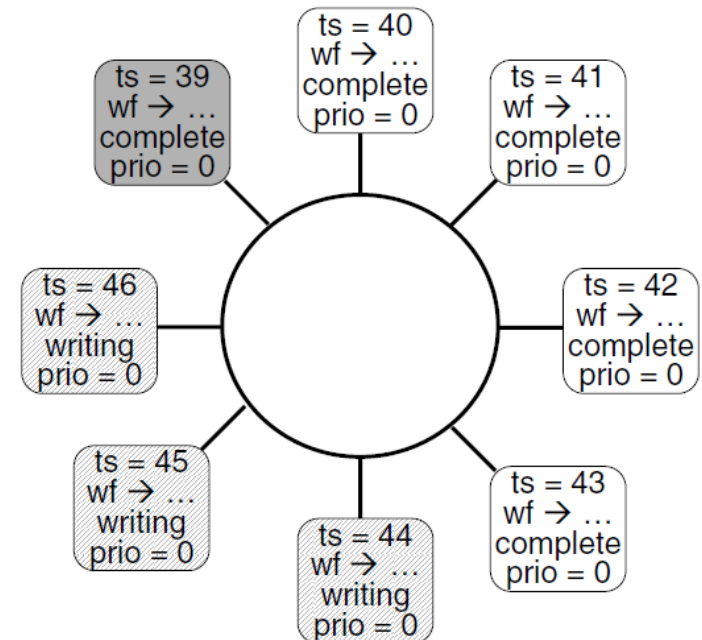


RingSTM

- `tx_begin()`
 - Its idea is to find oldest entry in the ring that is still writing back.
 - It depends on this invariant
$$L_i.st = \textit{writing} \implies \forall_{k>i} L_k.st = \textit{writing}$$
 - A transaction cannot change its status to complete if an older transaction is still writing
 - Guarantee detecting potential conflicts with the transactions still writing
 - Without waiting

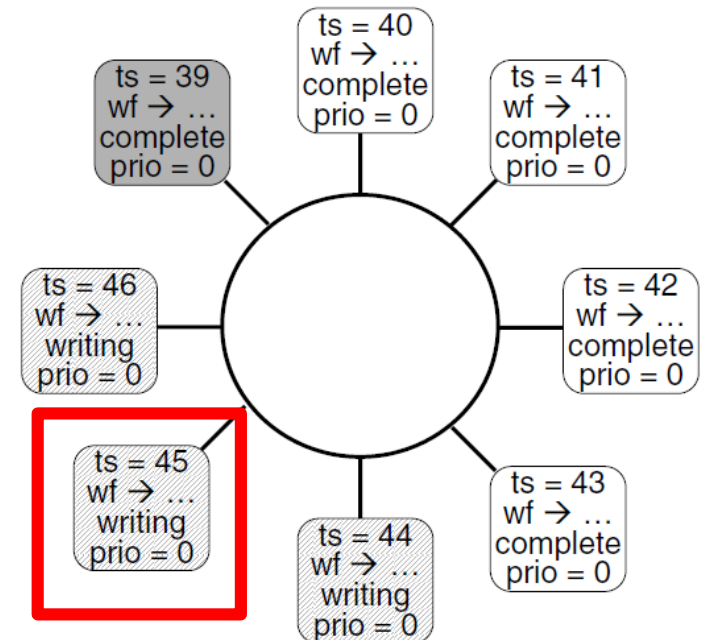
RingSTM

- `tx_begin()`
 - Its idea is to find oldest entry in the ring that is still writing back.
 - It depends on this invariant
$$L_i.st = \text{writing} \implies \forall_{k>i} l$$
 - A transaction cannot change older transaction is still writing
 - Guarantee detecting potent transactions still writing
 - Without waiting



RingSTM

- `tx_begin()`
 - Its idea is to find oldest entry in the ring that is still writing back.
 - It depends on this invariant
$$L_i.st = writing \implies \forall_{k>i} l$$
 - A transaction cannot change older transaction is still writing
 - Guarantee detecting potent transactions still writing
 - Without waiting



RingSTM

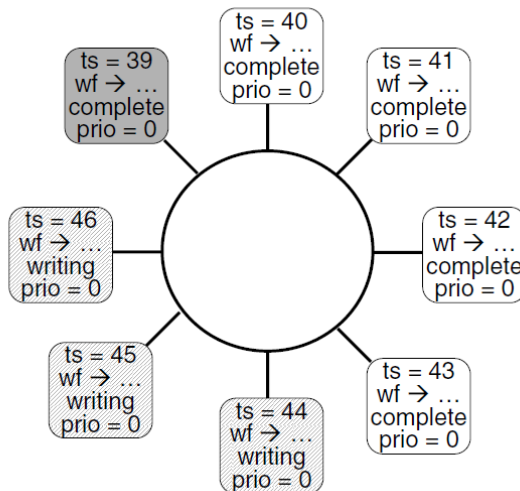
- tx_begin()
 - Reset thread-local metadata
 - $RV = \text{Global-Clock}$
 - while (ring[RV].status != complete ||
ring[RV].timestamp < RV)
 - RV--

We will assume infinite ring for simplicity
Note: Global-Clock is the current ring-index

RingSTM

- tx_begin()
 - Reset thread-local metadata
 - RV = Global-Clock
 - while (**ring[RV].status != complete** ||
ring[RV].timestamp < RV)
- RV--

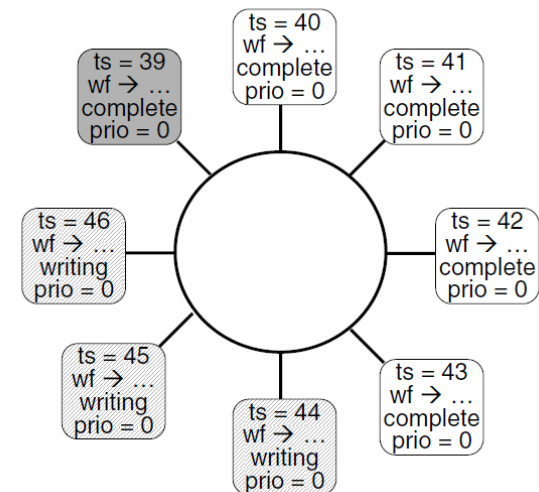
Lock backward for the first complete entry



RingSTM

- tx_begin()
 - Reset thread-local metadata
 - RV = Global-Clock
 - while (ring[RV].status != complete ||
ring[RV].timestamp < RV)
- RV--

The case when global-clock is incremented but the entry is not set yet



RingSTM

- tx_begin()
 - Reset thread-local metadata
 - RV = Global-Clock
 - while (ring[RV].status != complete ||
 ring[RV].timestamp < RV)
 - RV--

tm_begin:

```
1 read ring_index to TX.start
2 while ring[TX.start].st != complete ||
   ring[TX.start].ts < TX.start
3 TX.start--
4 fence(Read-Before-Read)
```

RingSTM

- `tx_write(addr, value)`
 - Add (or update) the `addr` and `value` to the write-set
 - Add the `addr` to the write-set signature

RingSTM

- `tx_write(addr, value)`
 - Add (or update) the `addr` and `value` to the write-set
 - Add the `addr` to the write-set signature

```
tm_write:  
1 TX.wset.add(addr, val)  
2 TX.wf.add(addr)
```

RingSTM

- tx_read(addr)
 - Find the addr is in the write-set signature
 - If found, find the addr is in the write-set
 - return the value buffered in the write-set
 - val = *addr
 - Add addr to read-set signature
 - tx_validate()
 - Return val

RingSTM

- tx_read(addr)
 - Find the addr is in the write-set signature
 - If found, find the addr is in the write-set
 - return the value buffered
 - val = *addr
 - Add addr to read-set signature
 - tx_validate()
 - Return val

Bloom filters lookup is very fast [O(1)], so it can be used to save the expensive write-set lookup
Note: TL2 uses a Bloom filter also

RingSTM

- tx_read(addr)
 - Find the addr is in the write-set signature
 - If found, find the addr is in the write-set
 - return the value buffered in the write-set
 - val = *addr
 - Add addr to read-set signature
 - tx_validate()
 - Return val

But, Bloom filter has false-positive
So, we have to confirm the address
is actually in the write-set

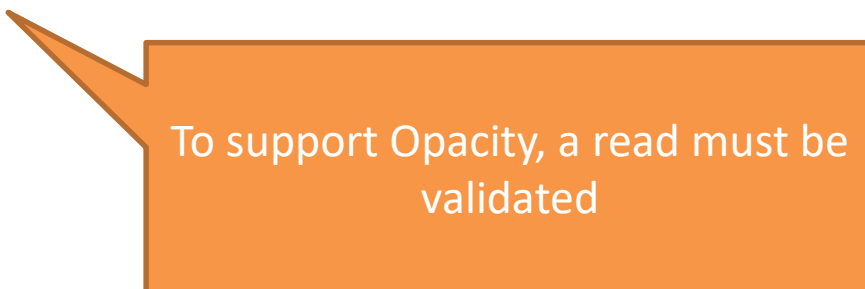
RingSTM

- tx_read(addr)
 - Find the addr is in the write-set signature
 - If found, find the addr is in the write-set
 - return the value buffered in the write-set
 - val = *addr
 - Add addr to read-set signature
 - tx_validate()
 - Return val

No read-set, only a read-set
signature

RingSTM

- `tx_read(addr)`
 - Find the `addr` is in the write-set signature
 - If found, find the `addr` is in the write-set
 - return the value buffered in the write-set
 - `val = *addr`
 - Add `addr` to read-set signature
 - `tx_validate()`
 - Return `val`



To support Opacity, a read must be validated

RingSTM

- tx_read(addr)

- Find the addr is in the write-set signature

- If found, find the addr is in the write-set

- return the value buffered in the write-set

- val = *addr

- Add addr to read-set signature

- tx_validate()

- Return val

```
tm_read:
```

```
1 if addr in TX.wf && addr in TX.wset
```

```
2   return lookup(addr, TX.wset)
```

```
3   val=*addr
```

```
4   TX.rf.add(addr)
```

```
5   fence(Read-Before-Read)
```

```
6   check()
```

```
7   return val
```

RingSTM

- tx_validate()
 - if Global-Clock == RV \rightarrow return
 - end = Global-Clock
 - while (ring[end].timestamp < end) wait
 - for ring entries between Global-Clock & RV+1
 - if (ring-entry.write-sig \cap read-set signature)
 - tx_abort()
 - if (ring-entry.status == writing)
 - end = (ring-entry-index) – 1
 - RV = end

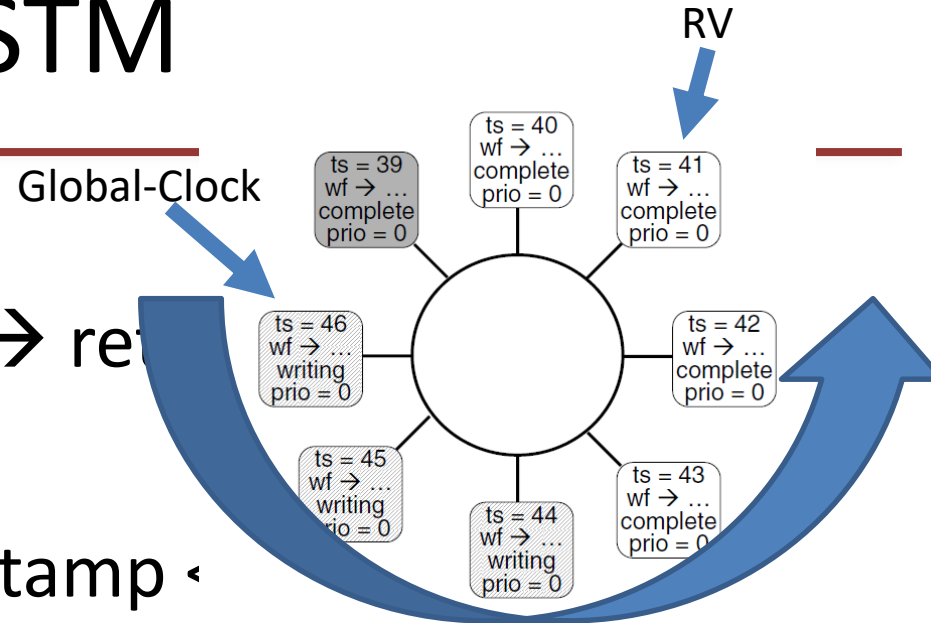
RingSTM

- tx_validate()
 - if Global-Clock == RV \rightarrow return
 - end = Global-Clock
 - while (ring[end].timestamp < end) wait
 - for ring entries between Global-Clock & RV+1
 - if (ring-entry.writing == true)
 - tx_abort()
 - if (ring-entry.status == writing)
 - end = (ring-entry-index) – 1
 - RV = end

The case when global-clock is incremented but the entry is not set yet

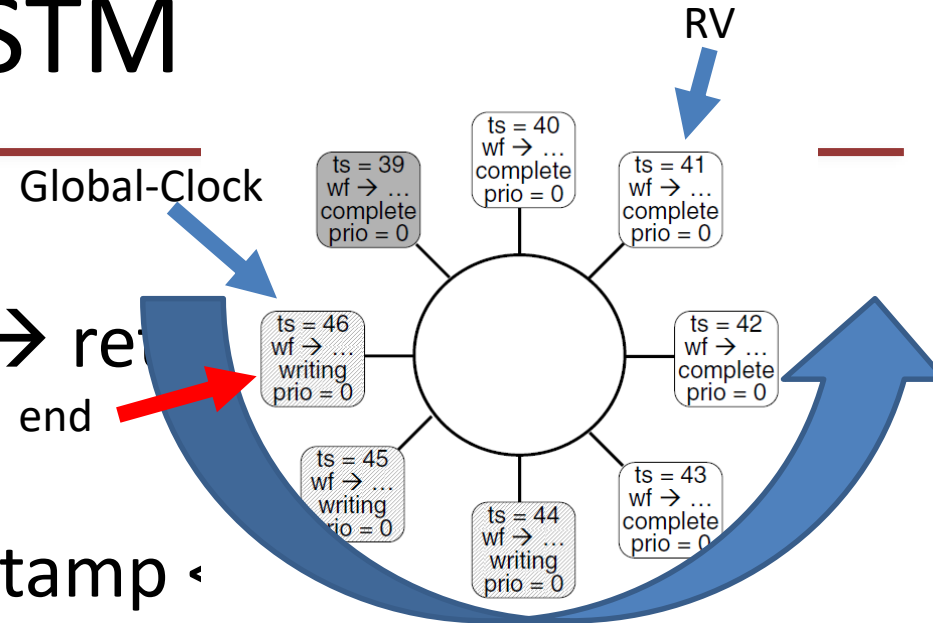
RingSTM

- tx_validate()
 - if Global-Clock == RV \rightarrow return
 - end = Global-Clock
 - while (ring[end].timestamp < Global-Clock)
 - for ring entries between Global-Clock & RV+1
 - if (ring-entry.write-sig \cap read-set signature)
 - tx_abort()
 - if (ring-entry.status == writing)
 - end = (ring-entry-index) – 1
 - RV = end



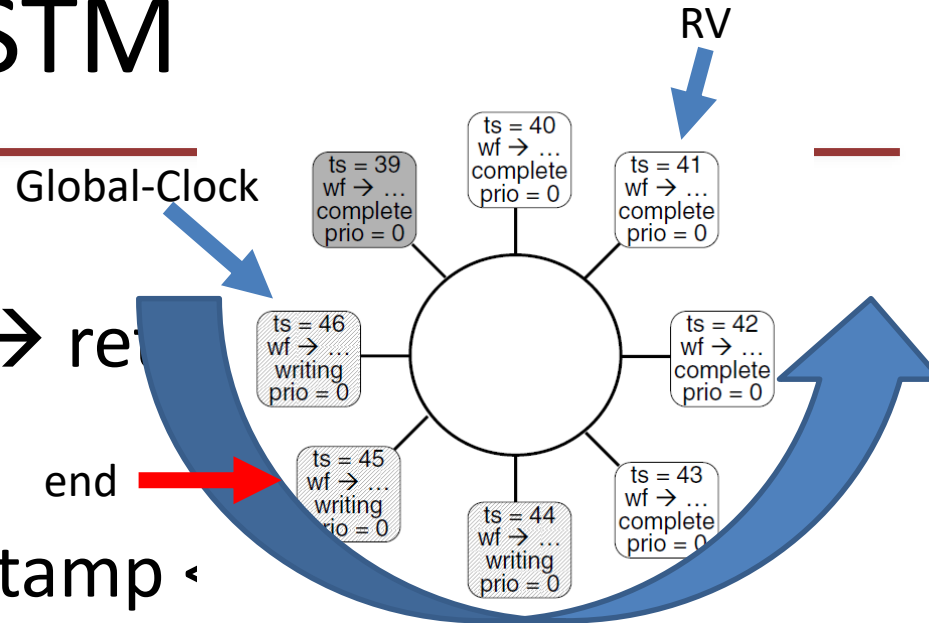
RingSTM

- `tx_validate()`
 - if `Global-Clock == RV` → `return`
 - `end = Global-Clock`
 - while (`ring[end].timestamp < Global-Clock`)
 - for ring entries between `Global-Clock` & `RV+1`
 - if (`ring-entry.write-sig ∩ read-set signature`)
 - `tx_abort()`
 - if (`ring-entry.status == writing`)
 - `end = (ring-entry-index) - 1`
 - `RV = end`



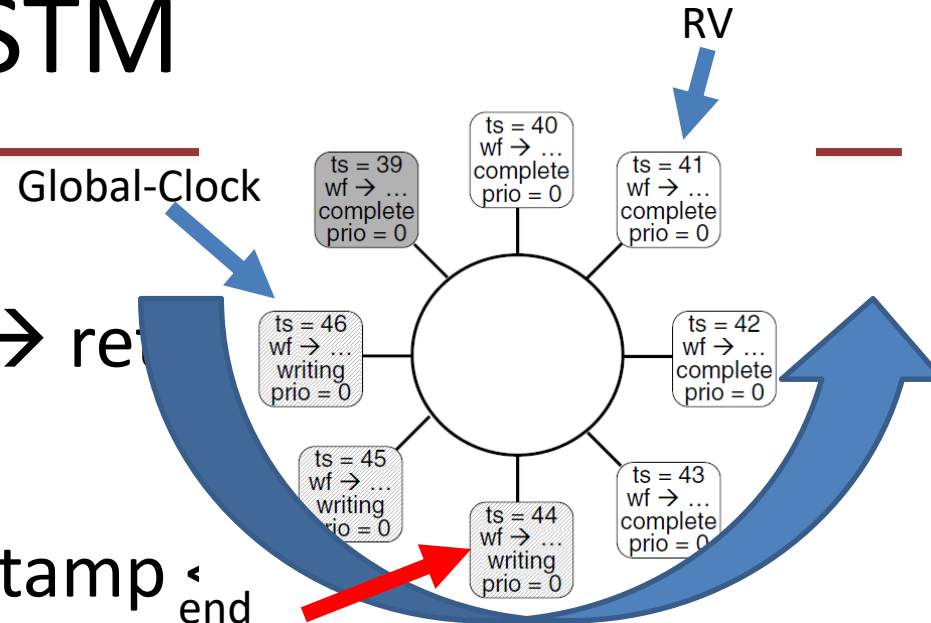
RingSTM

- `tx_validate()`
 - if `Global-Clock == RV` \rightarrow return
 - `end = Global-Clock`
 - while (`ring[end].timestamp <`
 - for ring entries between `Global-Clock` & `RV+1`
 - if (`ring-entry.write-sig` \cap `read-set signature`)
 - `tx_abort()`
 - if (`ring-entry.status == writing`)
 - `end = (ring-entry-index) - 1`
 - `RV = end`



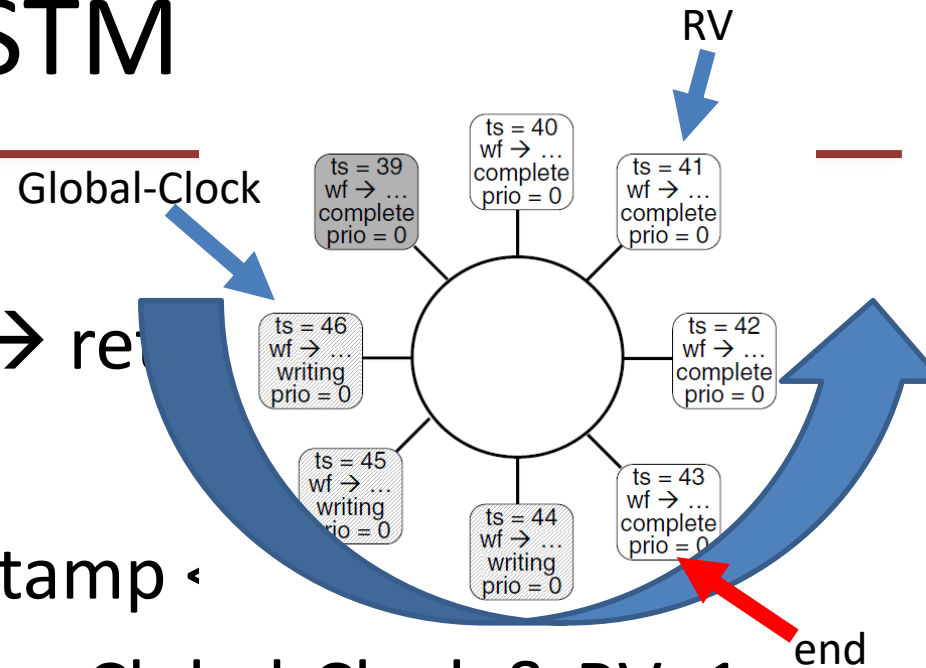
RingSTM

- `tx_validate()`
 - if `Global-Clock == RV` \rightarrow return
 - `end = Global-Clock`
 - while (`ring[end].timestamp` \leq `end`)
 - for ring entries between `Global-Clock` & `RV+1`
 - if (`ring-entry.write-sig` \cap `read-set signature`)
 - `tx_abort()`
 - if (`ring-entry.status == writing`)
 - `end = (ring-entry-index) - 1`
 - `RV = end`




RingSTM

- `tx_validate()`
 - if `Global-Clock == RV` \rightarrow return
 - `end = Global-Clock`
 - while (`ring[end].timestamp < Global-Clock`)
 - for ring entries between `Global-Clock` & `RV+1`
 - if (`ring-entry.write-sig` \cap `read-set signature`)
 - `tx_abort()`
 - if (`ring-entry.status == writing`)
 - `end = (ring-entry-index) - 1`
 - `RV = end`



RingSTM

- `tx_validate()`
 - if `Global-Clock == RV` → return
 - `end = Global-Clock`
 - while (`ring[end].timestamp < Global-Clock`)
 - for ring entries between `Global-Clock` & `RV+1`
 - if (`ring-entry.write-sig` \cap `read-set signature`)
 - `tx_abort()`
 - if (`ring-entry.status == writing`)
 - `end = (ring[end+1].timestamp > Global-Clock)`
 - **`RV = end`**



Timestamp extension. No need to check these entries again

RingSTM

check:

```
1 if ring_index==TX.start return
2 suffix_end=ring_index
3 while (ring[suffix_end].ts < suffix_end) SPIN
4 for i=ring_index downto TX.start+1
5   if intersect(ring[i].wf, TX.rf)
6     abort()
7   if ring[i].st==writing
8     suffix_end=i-1
9 TX.start=suffix_end
```


RingSTM

- tx_commit()
 - if (write-set.size == 0) → return // read-only tx
 - again: commit_time = Global-Clock
 - tx_validate()
 - If (!CAS(&Global-Clock, commit_time, commit_time + 1))
 - goto again
 - ring[commit_time + 1] = {writing, write-set-sig, commit_time + 1}
 - for (i = commit_time *downto* RV + 1)
 - if (ring[i].write-sig \cap write-set signature)
 - while(ring[i].status == writing) wait
 - For each entry in the write-set
 - *entry.addr = entry.value //Write back
 - while(ring[commit_time].status == writing) wait
 - ring[commit_time + 1].status = complete

RingSTM

- tx_commit()
 - if (write-set.size == 0) → return // read-only tx
 - again: commit_time = Global-Clock
 - tx_validate()
 - If (!CAS(&Global-Clock, commit_time, commit_time + 1))
 - goto again
 - ring[commit_time + 1] = {writing, ...}
 - for (i = commit_time *downto* RV + 1)
 - if (ring[i].write-sig \cap write-set signature)
 - while(ring[i].status == writing) wait
 - For each entry in the write-set
 - *entry.addr = entry.value //Write back
 - while(ring[commit_time].status == writing) wait
 - ring[commit_time + 1].status = complete

As in TL2 and NOrec, nothing extra is done for committing read-only transaction. It is linearized at the last tx_read's validation

RingSTM

- tx_commit()
 - if (write-set.size == 0) → return // read-only tx
 - again: commit_time = Global-Clock
 - tx_validate()
 - If (!CAS(&Global-Clock, +1))
 - goto again
 - ring[commit_time + 1] = {
 - for (i = commit_time do
 - if (ring[i].write-sig \cap write-set signature)
 - while(ring[i].status == writing) wait
 - For each entry in the write-set
 - *entry.addr = entry.value //Write back
 - while(ring[commit_time].status == writing) wait
 - ring[commit_time + 1].status = complete

Validation is done before adding an entry to the ring. This minimize contention window. Also, no need to add entries for transactions that will be aborted

RingSTM

- tx_commit()
 - if (write-set.size == 0) → return // read-only tx
 - again: commit_time = Global-Clock
 - tx_validate()
 - If (!CAS(&Global-Clock, commit_time, commit_time + 1))
 - goto again
 - ring[commit_time + 1] = {writing, ...}
 - for (i = commit_time *downto* 0)
 - if (ring[i].write-sig \cap write-set sig)
 - while(ring[i].status == writing) wait
 - For each entry in the write-set
 - *entry.addr = entry.value //Write back
 - while(ring[commit_time].status == writing) wait
 - ring[commit_time + 1].status = complete

If the transaction is valid, try to reserve an entry in the ring using the commit_time value (which is captured before the validation)

RingSTM

- tx_commit()
 - if (write-set.size == 0) → return // read-only tx
 - again: commit_time = Global-Clock
 - tx_validate()
 - If (!CAS(&Global-Clock, commit_time, commit_time + 1))
 - goto again
 - ring[commit_time].status = {writing, write-set-sig, commit_time + 1}
 - for (i = commit_time; i < commit_time + 1; i++)
 - if (ring[i].status == writing)
 - while (ring[i].status == writing) wait
 - For each entry in write-set
 - *entry.addr = entry.value // write back
 - while(ring[commit_time].status == writing) wait
 - ring[commit_time + 1].status = complete

If CAS failed, this means another transaction(s) has committed while doing the validation. This requires another validation (only for the newly committed transactions. Why?)

RingSTM

- tx_commit()
 - if (write-set.size == 0) → return // read-only tx
 - again: commit_time = Global-Clock
 - tx_validate()
 - If (!CAS(&Global-Clock, commit_time, commit_time + 1))
 - goto again
 - ring[commit_time + 1] = {writing, write-set-sig, commit_time + 1}
 - for (i = commit_time *downto* RV + 1)
 - if (ring[i].write-sig \cap write-set signature)
 - while(ring[i].status == writing) wait
 - For each entry in the write-set
 - *entry.addr = entry.value //Write back
 - while(ring[commit_time].status == writing) wait
 - ring[commit_time + 1].status = complete

At this point, my transaction is committed (it is valid and it has reserved its entry in the ring)

RingSTM

- tx_commit()
 - if (write-set.size == 0) → return // read-only tx
 - again: commit_time = Global-Clock
 - tx_validate()
 - If (!CAS(&Global-Clock, commit_time, commit_time + 1))
 - goto again
 - ring[commit_time + 1] = {writing, ...}
 - for (i = commit_time *downto* R)
 - if (ring[i].write-sig \cap write-set sig)
 - while(ring[i].status == writing) wait
 - For each entry in the write-set
 - *entry.addr = entry.value //Write back
 - while(ring[commit_time].status == writing) wait
 - ring[commit_time + 1].status = complete

This is the linearization point, a successful CAS

RingSTM

- tx_commit()
 - if (write-set.size == 0) → return // read-only tx
 - again: commit_time = Global-Clock
 - tx_validate()
 - If (!CAS(&Global-Clock, commit_time, commit_time + 1))
 - goto again
 - ring[commit_time + 1] = {writing, write-set-sig, commit_time + 1}
 - for (i = commit_time *downto* RV + 1)
 - if (ring[i].write-sig \cap write-set sig)
 - while(ring[i].status == writing) wait
 - For each entry in the write-set
 - *entry.addr = entry.value //Write back
 - while(ring[commit_time].status == writing) wait
 - ring[commit_time + 1].status =

Notice that the ring entry initially has $ts = 0$ (using our simplified infinite ring). At this line, the ring entry is populated. This is why we needed “ring[RV].timestamp < RV” in tx_begin and tx_validate

RingSTM

- tx_commit()
 - if (write-set.size == 0) → return // read-only tx
 - again: commit_time = Global-Clock
 - tx_validate()
 - If (!CAS(&Global-Clock, commit_time, commit_time + 1))
 - goto again
 - ring[commit_time + 1] = {writing, write-set-sig, commit_time + 1}
 - for (i = commit_time *downto* RV + 1)
 - if (ring[i].write-sig \cap write-set signature)
 - while(ring[i].status == writing) wait
 - For each entry in the write-set
 - *entry.addr = entry.value //Write back
 - while(ring[commit_time].status == writing)
 - ring[commit_time + 1].status = c

The write-set-sig acts as a write-lock. This preserve write-after-write ordering

RingSTM

- tx_commit()
 - if (write-set.size == 0) → return // read-only tx
 - again: commit_time = Global-Clock
 - tx_validate()
 - If (!CAS(&Global-Clock, commit_time, commit_time + 1))
 - goto again
 - ring[commit_time + 1] = {writing, write-set-sig, commit_time + 1}
 - for (i = commit_time *downto* RV + 1)
 - if (ring[i].write-sig \cap write-set signature)
 - while(ring[i].status == writing) wait
 - For each entry in the write-set
 - *entry.addr = entry.value //Write back
 - while(ring[commit_time].status == writing) wait
 - ring[commit_time + 1].status = complete

If there is no common elements between writing transactions, we can proceed with writing back in parallel.

RingSTM

- tx_commit()
 - if (write-set.size == 0) → return // read-only tx
 - again: commit_time = Global-Clock
 - tx_validate()
 - If (!CAS(&Global-Clock, commit_time, commit_time + 1))
 - goto again
 - ring[commit_time + 1] = {writing, write-set-sig, commit_time + 1}
 - for (i = commit_time *downto* RV)
 - if (ring[i].write-sig \cap write-set sig)
 - while(ring[i].status == writing) wait
 - For each entry in the write-set
 - *entry.addr = entry.value //Write back
 - while(ring[commit_time].status == writing) wait
 - ring[commit_time + 1].status = complete

Protect our invariant of not allowing a complete transaction until all previous transactions are complete

RingSTM

- tx_commit()
 - if (write-set.size == 0) → return // read-only tx
 - again: commit_time = Global-Clock
 - tx_validate()
 - If (!CAS(&Global-Clock, commit_time, commit_time + 1))
 - goto again
 - ring[commit_time + 1] = {writing, write-set-sig, commit_time + 1}
 - for (i = commit_time *downto* RV + 1)
 - if (ring[i].write-sig \cap write-set signature)
 - while(ring[i].status == writing) wait
 - For each entry in the write-set
 - *entry.addr = entry.value //Write back
 - while(ring[commit_time].status == writing) wait
 - ring[commit_time + 1].status = complete

Finally, mark the transaction's ring entry as complete which means all data are written to the memory and all previous entries in the ring are complete also

RingSTM

```
tm_end:
  1 if read-only return
  2 commit_time=ring_index
  3 check()
  4 if !CAS(ring_index, commit_time, commit_time+1)
  5   goto 2
  6 ring[commit_time+1]=(writing, TX.wf, commit_time+1)
  7 for i=commit_time downto TX.start+1
  8   if intersect(ring[i].wf, TX.wf)
  9     while ring[i].st=writing SPIN
  10 fence(ReadWrite-Before-Write))
  11 for (addr, val) in TX.wset
  12   write val at addr
  13 while ring[commit_time].st==writing SPIN
  14 fence(ReadWrite-Before-Write)
  15 ring[commit_time+1].st=complete
```

RingSTM

- tx_abort
 - //Just jump back to tx_begin to restart the transaction

RingSTM

- Pros
 - Lightweight read-set
 - Fast validation
 - Low memory overhead
 - Low cache pressure (one CAS operation)
 - Concurrent commits
 - Livelock freedom
- Cons
 - Imprecise validation
 - False conflict aborts
 - Sensitive to the Bloom filter size and hash function

Is RingSTM Opaque?

- Does it suffer from zombies?

T1:

- tx_begin
- tx_write(x, 2)
- tx_write(y, 4)
- tx_commit()

Initially $x=1$, $y=2$
Invariant: $y = 2x$

T2:

- tx_begin
- $v1 = \text{tx_read}(y)$
- $v2 = \text{tx_read}(x)$
- $\text{tx_write}(z, 1/(v1 - v2))$
- tx_commit()

Is RingSTM Opaque?

- Does it suffer from zombies?

T1:

- tx_begin
- tx_write(x, 2)
- tx_write(y, 4)
- tx_commit()

T2:

- tx_begin
- v1 = tx_read(y)
- v2 = tx_read(x)
- tx_write(z, 1/(v1 - v2))
- tx_commit()

NO zombies, it validates after each read. Any addition to the ring will trigger a validation after a tx_read. Y is in T2 read-set-sig and will intersect with T1 write-set-sig in the ring

RingSTM

T1:

- tx_begin
- v1 = tx_read(x)
- tx_write(y, v1 + 10)
- tx_commit()

T2:

- tx_begin
- v1 = tx_read(x)
- v2 = tx_read(y)
- tx_write(y, v1 + v2)
- tx_write(x, v1 + 1)
- tx_commit()

RingSTM

T1:

- tx_begin
- v1 = tx_read(x)
- tx_write(y, v1 + 10)
- tx_commit()

Will both transactions commit?
Are they Linearizable?

T2:

- tx_begin
- v1 = tx_read(x)
- v2 = tx_read(y)
- tx_write(y, v1 + v2)
- tx_write(x, v1 + 1)
- tx_commit()

RingSTM

T1:

- tx_begin
- v1 = tx_read(x)
- tx_write(y, v1 + 10)
- tx_commit()

An orange callout box with a speech bubble tail pointing towards the T1 transaction list.

IT DEPENDS!

T2:

- tx_begin
- v1 = tx_read(x)
- v2 = tx_read(y)
- tx_write(y, v1 + v2)
- tx_write(x, v1 + 1)
- tx_commit()

RingSTM

T1:

- tx_begin
- v1 = tx_read(x)
- tx_write(y, v1 + 10)
- tx_commit()

T2:

- tx_begin
- v1 = tx_read(x)
- v2 = tx_read(y)
- tx_write(y, v1 + v2)
- tx_write(x, v1 + 1)
- tx_commit()

We need to know the Bloom filter size and hash function.

Assume a Bloom filter of size 10 and a single hash function {addr%10}

Address of x is 10

Address of y is 100

RingSTM

T1:

- tx_begin
- v1 = tx_read(x)
- tx_write(y, v1 + 10)
 - (Addr of y) % 10 \rightarrow 100 % 10 = 0
- tx_commit()
 - Write-set-sig = [1,0,0,0,0,0,0,0,0,0]
 - Global-Clock = 1

We need to know the Bloom filter size and hash function.

Assume a Bloom filter of size 10 and a single hash function {addr%10}

Address of x is 10

Address of y is 100

T2:

- tx_begin (RV = 0)
- v1 = tx_read(x)
 - (Addr of x) % 10 \rightarrow 10 % 10 = 0
 - Read-set-sig = [1,0,0,0,0,0,0,0,0,0]
- v2 = tx_read(y)
 - tx_validate \rightarrow
[1,0,0,0,0,0,0,0,0,0] \cap [1,0,0,0,0,0,0,0,0,0] \neq 0
 \rightarrow tx_abort()
- tx_write(y, v1 + v2)
- tx_write(x, v1 + 1)
- tx_commit()

RingSTM

T1:

- tx_begin
- v1 = tx_read(x)
- tx_write(y, v1 + 10)
 - (Addr of y) % 10 \rightarrow 100 % 10 = 0
- tx_commit()
 - Write-set-sig = [1,0,0,0,0,0,0,0,0,0]
 - Global-Clock = 1

NOW

Assume a Bloom filter of size 10 and a single hash function {addr%10}

Address of x is 1

Address of y is 100

T2:

- tx_begin (RV = 0)
- v1 = tx_read(x)
 - (Addr of x) % 10 \rightarrow 1 % 10 = 1
 - Read-set-sig = [0,1,0,0,0,0,0,0,0,0]
- v2 = tx_read(y)
 - tx_validate \rightarrow
[1,0,0,0,0,0,0,0,0,0] \cap [0,1,0,0,0,0,0,0,0,0] == 0
 \rightarrow Valid \rightarrow RV = 1
- tx_write(y, v1 + v2)
- tx_write(x, v1 + 1)
- tx_commit()

RingSTM

T1:

- tx_begin
- v1 = tx_read(x)
- tx_write(y, v1 + 10)
 - (Addr of y) % 10 \rightarrow 100 % 10 = 0
- tx_commit()
 - Write-set-sig = [1,0,0,0,0,0,0,0,0,0]
 - Global-Clock = 1

With RingSTM, there can be false conflicts which depends on how the Bloom filter is implemented

T2:

- tx_begin (RV = 0)
- v1 = tx_read(x)
 - (Addr of x) % 10 \rightarrow 1 % 10 = 1
 - Read-set-sig = [0,1,0,0,0,0,0,0,0,0]
- v2 = tx_read(y)
 - tx_validate \rightarrow
[1,0,0,0,0,0,0,0,0,0] \cap [0,1,0,0,0,0,0,0,0,0] == 0
 \rightarrow Valid \rightarrow RV = 1
- tx_write(y, v1 + v2)
- tx_write(x, v1 + 1)
- tx_commit()

RingSTM

T1:

- tx_begin
- v1 = tx_read(x)
- tx_write(y, v1 + 10)
- tx_commit()

T2:

- tx_begin
- v1 = tx_read(x)
- v2 = tx_read(y)
- tx_write(y, v1 + v2)
- tx_write(x, v1 + 1)
- tx_commit()

RingSTM

T1:

- tx_begin
- v1 = tx_read(x)
- tx_write(y, v1 + 10)
- tx_commit()

T2:

- tx_begin
- v1 = tx_read(x)
- v2 = tx_read(y)
- tx_write(y, v1 + v2)
- tx_write(x, v1 + 1)
- tx_commit()

Will both transactions commit?
Are they Linearizable?

RingSTM

T1:

- tx_begin
- tx_write(x, 15)
- v1 = tx_read(x)
- tx_write(y, v1 + 15)
- tx_commit()

T2:

- tx_begin
- v1 = tx_read(x)
- v2 = tx_read(y)
- tx_write(y, v1 + v2)
- tx_write(x, v1 + 1)
- tx_commit()

RingSTM

T1:

- tx_begin
- tx_write(x, 15)
- v1 = tx_read(x)
- tx_write(y, v1 + 15)
- tx_commit()

T2:

- tx_begin
- v1 = tx_read(x)
- v2 = tx_read(y)
- tx_write(y, v1 + v2)
- tx_write(x, v1 + 1)
- tx_commit()

What about this scenario?

RingSTM

T1:

- tx_begin
- v1 = tx_read(x)
- tx_write(y, v1 + 15)
- tx_commit()

T2:

- tx_begin
- v1 = tx_read(x)
- v2 = tx_read(y)
- tx_write(y, v1 + v2)
- tx_write(x, v1 + 1)
- tx_commit()

RingSTM

T1:

- tx_begin
- v1 = tx_read(x)
- tx_write(y, v1 + 15)
- tx_commit()

T2:

- tx_begin
- v1 = tx_read(x)
- v2 = tx_read(y)
- tx_write(y, v1 + v2)
- tx_write(x, v1 + 1)
- tx_commit()



What about this scenario?

RingSTM

T1:

- tx_begin

T2:

- tx_begin

read(x)

read(y)

+ v2)

+ 1)

tx_commit()

In the previous 3 scenarios, Bloom filter implementation will not affect! WHY?

- v1 = tx_read(x)
- tx_write(y, v1 + 15)
- tx_commit()